

Intro to Computational Complexity Theory

Suhas Arehalli

COMP221 - Spring 2024

1 Intro to Computational Complexity

- So far, we've been worried about how fast particular algorithms are.
 - But often, we want to know how fast *any* solution to that problem can be.
 - i.e., “Could I write a faster algorithm to solve this problem?”
- In general, it's hard to prove that no faster solution is possible!
 - These are often trivial or near-impossible!
 - i.e., search for unsorted arrays can't be faster than linear time, since we must at least check each index.
 - A more sophisticated argument of this kind is the problem from HW2 where you proved *all* comparison-based sorts are $\Omega(n \log n)$!
- Instead, we'll adopt a *relative* strategy; we will prove a particular problem is at least as hard as some other canonical problem that's known to be difficult!

1.1 (Turing/Cook) Reductions

- Consider the following scenario:
 - We have some problem specification X that we want to prove is hard (i.e., doesn't have a fast solution).
 - Assume we also have some problem specification Y that we *know* is hard — i.e., we know *every* solution to Problem Y is in $\Omega(h(n))$.
 - Let PROBX be an arbitrary algorithm that solves problem X that runs in worst-case time complexity $\Theta(f(n))$
 - Now consider a function TRANSLATE that turns a particular *solution* of problem X into a solution for problem Y in time $\Theta(g(n))$.
 - We can then construct a solution to Problem Y in terms of PROBX:

```
function PROBY(...)
   $x \leftarrow$  PROBX(...)
  return TRANSLATE( $x$ )
```

end function

- Consider the time complexity of this algorithm: $\Theta(f(n) + g(n))!$
- Since this is a solution to Problem Y, and we know Problem Y is hard, we know this solution is $\Omega(h(n))$.
- Thus we know that $f(n) + g(n) \in h(n)!$
- In the best case, $g(n)$ is dominated by $f(n)/h(n)$, and we can conclude that $f(n) \in \Omega(h(n))!$
- Let's sit with that for a second: If we find a sufficiently fast translation from a solution to Problem X to a solution to Problem Y, **and** we know Problem Y is sufficiently hard, then we know that Problem X is also hard!
 - Hard, so far, means that any algorithm that solves the problem is slower than some bound. We'll hone in on what that bound is in a bit!
- Note that I've elided translation of the *inputs* of the problem — this is actually perhaps the *most* interesting part of the translation process, but bear with me for now!
- For now, we're interested in writing functions like PROBY that use calls to algorithms that solve Problem X (i.e., PROBX) in order to solve Problem Y.
 - This is called a *reduction* of Problem Y to Problem X.
 - Note the ordering here! It's tricky!
 - * *Reducing* Problem Y to Problem X means *translating* a solution (or multiple solutions) to an instance of Problem X into a solution to Problem Y.
 - * A fast reduction of Problem Y to Problem X proves that Problem X is at least as hard as Problem Y.
 - More technically, the general class of reductions we're considering right now are called *Turing* or *Cook Reductions*. We'll look at a different, more specific kind of reduction later.

1.2 Decision Problems

- In general, it's messy to trying to work with the various kinds of problem specifications we've seen so far.
- In the future, we will discuss complexity theory in terms of *Decision Problems* — problems with yes/no (i.e., true or false) answers.
- We do this because of these kinds of problems are well studied and have a lot of theory around them (i.e., from automata/formal language theory!).
 - To fully understand what's going on here, you should take Theory of Computation!
 - For now, you'll see this pay dividends when we talk about Karp/Many-One reductions, and in general when these problems are just easier to work with.

- The claim underlying this choice is that we can convert interesting algorithmic problem specifications into equally interesting (i.e., just as hard!) decision problem variants!
- To see this, consider the case of *Vertex Cover*:
 - A *Vertex Cover* for a Graph $G = (V, E)$ is a set $C \subseteq V$ such that $\forall e \in E$, some vertex $v \in C$ is incident to e . That is, $\forall (v, w) \in E$, $v \in C$ or $w \in C$ (or both!).
 - **(Minimum) Vertex Cover**: Given a Graph $G = (V, E)$, find the *smallest* vertex cover $C \subseteq V$.
 - **Vertex Cover (Decision)**: Given a Graph $G = (V, E)$ and $k \in \mathbb{Z}_{\geq 0}$, determine whether there is a vertex cover C with $|C| \leq k$
- My claim: We can reduce (minimum) Vertex Cover to its decision variant and vice-versa.
- Consider the following reduction:

```

function VERTEXCOVER_DECISION( $G = (V, E)$ ,  $k$ )
   $C \leftarrow \text{VERTEXCOVER}(G)$ 
  if  $|C| \leq k$  then
    return true
  end if
  return false
end function

```

- This reduces the decision problem to the optimization problem with a constant time transformation of the output!
 - This means that if VERTEXCOVER runs in $\Theta(f(|V|))$, then VERTEXCOVER_DECISION can be solved in $\Theta(f(|V|))$ as well!
 - That is, we showed that VERTEXCOVER is at least as hard as VERTEXCOVER_DECISION!
- Of course, this is not surprising: Most of you probably expected that the decision variant was easier than the optimization version!
- So let's consider the trickier reduction:¹

```

function VERTEXCOVER( $G = (V, E)$ )
   $k \leftarrow 0$ 
  while  $\neg \text{VERTEXCOVER\_DECISION}(G, k)$  do
     $k \leftarrow k + 1$ 
  end while
   $C \leftarrow \emptyset$ 
  for  $v \in V$  do
     $E' \leftarrow E \setminus \{e \in E \mid v \text{ incident to } e\}$ 
     $V' \leftarrow V \setminus \{v\}$ 
     $G' \leftarrow (V', E')$ 
    if VERTEXCOVER_DECISION( $G', k - 1$ ) then

```

¹This is different than the algorithm I gave you in class — this one is more straightforward, and doesn't require some of the handwaving I did about modifying the edge set!

```

     $C \leftarrow C \cup \{v\}$ 
     $E \leftarrow E'$ 
     $V \leftarrow V'$ 
     $k \leftarrow k - 1$ 
    if  $k == 0$  then
        return  $C$ 
    end if
end if
end for
end function

```

- To see that this algorithm is correct, observe...
 - ...that the first for-loop will set k to the size of the minimum vertex cover
 - ...that if we enter the if-statement, the minimum vertex cover of size $k - 1$ for G' (call it C') covers all edges that v is no incident to!
 - Thus, $C' \cup \{v\}$ is a vertex cover of G of size k !
 - i.e., at each iteration, $C \subseteq C^*$, where C^* is a minimal vertex cover of (initial) G .
 - Write out a formal proof as practice!
- Assume that VERTEXCOVER_DECISION runs in $\Theta(f(|V|))$ time. Then this algorithm runs in $\Theta(|V|f(|V|))$ time!
- Is this good enough?
 - This is a *weird* reduction — we call the decision problem's algorithm order $|V|$ times!
 - In terms of asymptotic time complexity, this reduction is an order of magnitude slower!
 - I'll still claim that this difference doesn't matter *for the kinds of things we care about*!
- What do we care about? *Tractability*!
- A problem X is *tractable* if there exists an algorithm that solves that problem in polynomial time. That is, for some $c \in \mathbb{Z}$, there exists a $O(n^c)$ algorithm that solves problem X !
 - Note that this means that an algorithm is *intractable* if there exist **no** polynomial time solutions.
 - Proving tractability means finding a polynomial time algorithm to solve X . Proving intractability means proving **all** algorithms that solve problem X run slower than polynomial time!
- Note that, given this second reduction, we can still claim that if VERTEXCOVER_DECISION is tractable, then VERTEXCOVER is tractable!
 - Perhaps more importantly, this reduction tells us that if VERTEXCOVER is *intractable*, then VERTEXCOVER_DECISION is also intractable!
- With both reductions, we know that, by our tractability-based definition of hardness, both problems are equally hard!
- Thus, moving forward, we'll be working nearly exclusively with decision problems rather than more complex problem specifications.

2 (Karp) Reductions

- Now we will stay squarely in the world of decision problems, and this buys us the possibility of a stronger (and more theoretically satisfying!) kind of reduction: A *Karp*, or *Many-to-one* reduction!
- As promised, this is also where we start thinking about translating from inputs to inputs!
- Let X and Y be problems, PROB_X and PROB_Y algorithms that solve those problems. A *Karp* reduction from Y to X is a transformation TRANSLATE such that for any input I to PROB_Y , $\text{PROB}_Y(I) = \text{PROB}_X(\text{TRANSLATE}(I))$.

– That is, in algorithmic form, we can write a solution to Problem Y in the form

```
function  $\text{PROB}_Y(I)$ 
     $I' \leftarrow \text{TRANSLATE}(I)$ 
    return  $\text{PROB}_X(I')$ 
end function
```

- This is a very powerful kind of reduction: we know that Problem Y isn't just solvable using Problem X (like in a Turing Reduction), but that instances of Problem Y are *special cases* of Problem X !
 - TRANSLATE then tells us exactly which instances of Problem X correspond to each instance of Problem Y .
- Two things are critical here: Because we're translating problem instances/inputs,
 - ... we can only call an algorithm that solves problem X exactly once!
 - ... we can't transform the *output* of the algorithm in any way!
- The first point is a big restriction, but the second is subtly important.
 - Constructing these kinds of reductions is only really possible if we work with decision problems, because then we just need both algorithms to say true or false!
 - This also means that even within decision problems, we can't do things like negate outputs — We have to construct a mapping such that instance I of problem X is true iff some instance I' of problem Y is true!
- From now forward, when we discuss reductions for decision problems, I'm going to be focusing on Karp reductions.
 - If any of you are interested in Theoretical CS, this is the kind of reduction you'll see by default (Skiena doesn't even draw this distinction).

3 Canonically Hard Problems

- Remember our goal here is to convince ourselves some problems are hard.
 - We now have a theory that helps us say Problem X is at least as hard as Problem Y (i.e., we can reduce Problem Y to Problem X).

- To show some Problem X is hard, we now need to find sufficiently hard problems Y that we can reduce to X!
- Remember our formalization of the idea of a problem being hard is a problem being *intractable*.
- It follows that our goal should be to find intractable problems to reduce to problems we think are hard!
- If we want to be really careful about the term intractable, we’re kind of out of luck...
 - What if we’re just not clever enough to find a polynomial time solution?
- Instead, we will “settle” for problems that we’re very confident are intractable: problems that are *NP-hard*
- A problem in complexity class *NP* (Non-deterministic Polynomial time) is able to be *verified true* in polynomial time.
 - A *verifier* is an algorithm that, for a particular problem, takes in a problem instance and a *certificate/witness* — some information that serves as evidence that the purported answer is correct — and concludes whether the witness justifies a true or false answer.
 - There is some subtlety here: What we want is that there *exists* some certificate such that our verifier will confirm that the right answer is correct in polynomial time.
 - This witness essentially acts as a hint to our algorithm: Is there some information that I can give you such that you can solve the problem in polynomial time?
 - * In practice, that hint is usually in the form of the solution to the non-decision version of the problem: A vertex cover, or a variable assignment, or Hamiltonian path, etc.
 - For instance, consider the HAMILTONIANCYCLE decision problem. A *certificate* would be the sequence of edges in the graph, and a *verifier* for true answers would be an algorithm like

```

function HP_VERIFIER( $G = (V, E), C$ )
   $(v, w) \leftarrow C[1]$ 
   $curr \leftarrow v$ 
  for  $(v, w) \in C$  do
    if  $(v, w) \notin E$  or  $v \neq curr$  or  $w$  is visited then
      return false
    end if
     $curr \leftarrow w$ 
  end for
  for  $v \in V$  do
    if  $v$  not visited then
      return false
    end if
  end for
  return true
end function

```

- For any case where the Hamiltonian Path problem’s answer is true, we can *verify* that answer by providing a certificate: a simple cycle in G that visits every vertex!
- For those who want to go above and beyond: There is a complementary class of problems called $coNP$, which is the set of problems that can be verified *false* in polynomial time. Whether $NP = coNP$ is an open question!
- The relationship between this verification definition and the *nondeterministic* in NP comes from automata theory. You can think about nondeterminism as a sort of parallelism: If you could try to verify every possible certificate at once, you can solve the problem in polynomial time!
- It should be intuitive that $P \subseteq NP$: If you can solve it from scratch in polynomial time, you can verify the problem in polynomial time.
- However, we’re not certain polynomial time solutions aren’t possible for problems in NP !
 - Literally a million dollar problem! See the Millenium Prize Problems!
- But, given all the attempts at solving these problems, and the fact that a large group of these problems have been reduced to each other (NP -*Complete* problems), they are hard enough for our purposes
 - *Functionally* intractable!

3.1 SAT/3-SAT

First, some terms:

- A *variable* v takes boolean values
- A *literal* is either a variable v or its negation $\neg v$
- A *clause* is a chain of disjunctions (logical ors) of literals: $v_1 \vee v_2 \vee \neg v_3$

Then we get *SAT* and *3-SAT*!

SAT

For: A set of variables V and a set C of clauses constructed from those variables.

Determine: Whether there exist a set of assignments of true or false to each variable such that each of the clauses evaluates to true.

3-SAT

For: A set of variables V and a set C of clauses of 3 *literals* constructed from those variables.

Determine: Whether there exist a set of assignments of true or false to each variable such that each of the clauses evaluates to true.

3.2 Hamiltonian Cycle/Path

The unweighted version of the traveling salesman problem

Hamiltonian Cycle/Path

For: A Graph $G = (V, E)$

Determine: Whether there exists a cycle/path that visits every vertex in V exactly once.

3.3 Vertex Cover/ Independent Set

Vertex Cover and Vertex Cover's counterpart!

A *vertex cover* for an (undirected) graph $G = (V, E)$ is a set $C \subseteq V$ such that for every $(v, w) \in E$, either $v \in C$ or $w \in C$.

An *independent set* for an (undirected) graph $G = (V, E)$ is a set $S \subseteq V$ such that for any $v, w \in S$, $(v, w) \notin E$

Vertex Cover
For: A Graph $G = (V, E)$ and $k \in \mathbb{Z}_{\geq 0}$
Determine: Whether there exists a vertex cover C with $ C \leq k$.
Independent Set
For: A Graph $G = (V, E)$ and $k \in \mathbb{Z}_{\geq 0}$
Determine: Whether there exists an independent set S with $ S \geq k$.

3.4 Integer Partition

A *multi-set* is just a set that allows duplicates. That is, an un-ordered collection with repetitions.

Integer Partition
For: A multi-set of integers S
Determine: If there exists some partition S_1, S_2 of S such that $\sum_{s \in S_1} s = \sum_{s \in S_2} s$

3.5 And others

It's probably useful to be familiar with other, related NP-Complete problems: Clique shows up often, or something like Linear Integer Programming is probably of practical interest. But the few here (Vertex Cover/Independent Set, SAT/3-SAT, Hamiltonian Cycle/Path, and Integer Partition) cover a few general classes of problems: vertex selection in graphs (Vertex Cover/Independent Set), edge selection in graphs (Hamiltonian Cycle/Path), logic problems (SAT/3-SAT), and combinatorics/counting problems (Integer Partition). As Skiena suggests, and at least for this class, these will cover our bases!

4 Reductions to show hardness

- Since $P \subseteq NP$, reducing to a problem X from a problem in NP is not quite enough to show that the problem is practically intractable.
- Instead, we need to be concerned with problems that are at least as hard as everything in NP (i.e., it's one of the hardest problems in NP). These problems are handily called *NP-Hard*
 - Luckily we've developed a theory that lets us answer this question: A problem is NP-Hard iff every problem in NP can be reduced (in polynomial time) to that problem.
 - Note that this doesn't mean that NP-Hard problems have to be in NP ! They can be harder — impossible to verify correct in polynomial time!
 - Problems that are NP-Hard **and** are in NP are called *NP-Complete*

- To be fully clear: What this means is if *any* NP-hard problem could be solved in polynomial time, *every* NP problem could be, using the polynomial time reductions. The fact that no one's solved any of the NP-Complete problems would make it really surprising if $P = NP$!
- The canonical problems in the previous section are all NP-Complete, so they make a good target for reductions!

Example: Consider The Traveling Salesman Problem:

Traveling Salesperson Problem

For: A *complete* graph $G = (V, E)$, a cost function $c : E \rightarrow R$, and a bound k

Determine: Whether a cycle that visits every vertex exactly once exists with total cost $\leq k$.

To show that the traveling salesman problem is *NP – Hard*, we simply need to write a Karp reduction that solves an NP-Complete problem from the previous section using a solution to the TSP.

Now, thinking through our canonical problems, the TSP sounds a lot like a Hamiltonian Cycle, but for a weighted graph! So we can construct a transformation that solves the Hamiltonian Path problem in terms of the TSP. Note that in the standard formulation, it's assumed that the given G for the TSP is *complete*, so we will need to modify our graph.

Consider

```

function HAMILTONIANPATH( $G = (V, E)$ )
   $E' \leftarrow V^2$ 
  for  $(v, w) \in E'$  do
    if  $(v, w) \in E$  then
       $c'((v, w)) \leftarrow 1$ 
    else
       $c'((v, w)) \leftarrow \infty$ 
    end if
  end for
  return TSP( $(V, E'), c', |V|$ )
end function

```

This will return *true* iff a a tour that visits every vertex exists in $G' = (V, E')$ with cost function c' , which, given c' , can only exist if the circuit uses only the edges that are also in E . Thus, that TSP solution is also a hamiltonian path in G , and thus this algorithm return true iff there is a Hamiltonian path in G . Of course, this solution also runs in polynomial time: Everything else in the transformation is dominated by the $\Theta(|V|^2)$ time to construct $c' : E' \rightarrow R$.

Thus the TSP is NP-Hard! To show that it's NP-Complete, we'd simply need to construct a verifier that runs in polynomial time (A good exercise!).

5 “Solving” NP-Hard Problems Fast

- Proving a problem is (NP-)hard typically doesn't change the fact that the problem needs to be solved somehow.

- The practical question we have to ask is how we can solve those problems *as fast as we can*.
- There are a few approaches:
 1. Speed up the non-polynomial solution:
 - This usually involves techniques based on pruning, or trying to search in an order that will reach the solution before exhausting the full search space.
 - Think something like Best-First Search or A*. Sometimes you can prune branches from a search tree!
 - There may also be algorithms with exponential or factorial time complexities in the worst case, but have polynomial time average case performance!
 2. You can also run polynomial time algorithms that approximate the correct solution.
 - In these cases, we often want to be able to estimate *how far* we are from an optimal solution
 - We often write an algorithm is a k -approximation if it's within a factor of k from the optimal solution
 - * That is, A k approximation of an TSP algorithm constructs a tour with a cost at worst k times larger than the optimal TSP tour.
 3. Sometimes, we can write straightforward *probabilistic* polynomial time solutions that, in expectation (i.e., on average) produce fairly good approximations of the true solution.

Example: Consider the Vertex Cover problem. We must construct something close to the minimum vertex cover of the graph, so if we incrementally add vertices to a partial cover C , we want to minimize the number of vertices we add. However, because it needs to be a cover, every edge needs at least once of its vertices in C .

So here's our approach: We iterate through edges, and if they're uncovered, we know one of the two edges must be in the minimum vertex cover C^* . We don't have a clever approach to know which vertex is in C^* , so we'll guess at random!

```

function RANDOMVERTEXCOVER( $G = (V, E)$ )
   $C \leftarrow \emptyset$ 
  for  $(v, w) \in E$  do
    if  $v \notin C$  and  $w \notin C$  then
       $x \leftarrow \text{RANDOMSELECT}(v, w)$ 
       $C \leftarrow C \cup \{x\}$ 
    end if
  end for
end function

```

My claim is that this is, in expectation, a 2-approximation of VERTEXCOVER.

Proof: Consider the true minimum vertex cover of G , C^* : For each edge $e \in E$, by the definition of a vertex cover, $\exists v^* \in C^*$ that is incident to e . Note that this means for each edge we iterate through in our for-loop, we have a 50% chance of picking a vertex $v^* \in C^*$, and if we pick the other vertex, we will at worst have to reconsider v^* at another edge incident to v^* , and there are at most $\text{degree}(v^*)$ of those. Thus, for each vertex v^*

we expect to add less than

$$\sum_{i=1}^{\text{degree}(v^*)} \frac{i}{2^i}$$

vertices. That is, there is a 50% chance we get it right on the first edge (and we add just v^*). If we get it wrong the first time, there's a 50% chance we get it right the second time, which gives a $\frac{1}{2} \cdot \frac{1}{2} = \frac{1}{4}$ probability that we got it wrong first and right second, and end up with 2 vertices added to C : v^* and some other vertex. Continuing this process, we can generalize to say that there is a $\frac{1}{2^i}$ probability that we add at most i vertices to cover the edges v^* covered in the minimum vertex cover. Thus, the expected upper bound of the number vertices we add to C to cover what v^* covered is $\sum \frac{i}{2^i}$.

Observe that this is an upper bound because we add vertices until we select v^* , which is guaranteed to cover all of the edges it covers in the minimum vertex cover. It's possible that we don't actually need $v^* \in C$ to form a vertex cover, but this upper bound will be sufficient for us.

Some clever tricks can show that $\sum_{i=1}^{\infty} \frac{i}{2^i} = 2^2$, and since all the terms of the sum are positive, $\sum_{i=1}^{\text{degree}(v^*)} \frac{i}{2^i} \leq 2$. Thus, for every $v^* \in C^*$, we expect to add fewer than 2 vertices to C . Thus, $\mathbb{E}[|C|] \leq 2|C^*|$ (i.e., we expect $|C| \leq 2|C^*|$).

Note that the probability theory (and identities from calculus) necessary for these probabilistic analyses are no expected background for this course, so like QuickSort analyses earlier in the class, you will not be expected to reproduce these analysis independently. If if ever ask anything related to average-case or expected error analyses, I'll provide the relevant results from probability!

²Let's use some calculus: Let $f(n) = \sum_{i=0}^{\infty} nfty n^i$. We know that this is an infinite geometric series, and thus we also know $f(n) = (1 - n)^{-1}$. Now observe that the derivative $f'(n) = \sum_{i=0}^{\infty} nfty(i)n^{i-1}$ from the sum version, and $f'(n) = (1 - n)^{-2}$ from the closed form geometric series solution. Then observe that $\frac{1}{2}f'(\frac{1}{2})$, using the sum version, is exactly what we're looking for! So we have

$$\begin{aligned} nf'(n) &= \sum_{i=1}^{\infty} in^i = n(1 - n)^{-2} \\ \frac{1}{2}f'(\frac{1}{2}) &= \sum_{i=1}^{\infty} \frac{i}{2^i} = \frac{1}{2}(1 - \frac{1}{2})^{-2} \\ &= \frac{1}{2} \cdot 4 \\ &= 2 \end{aligned}$$